

Incident Investigation Playbook (Reference)

Sample reference document • For OpsPilot agent grounding/evaluation • Fictional environment for demo purposes

1. Bad Deployment

Recognize it by: latency or error-rate increase that starts shortly after a deployment; the new version shows a much higher error rate than the previous one.

Typical path: check the service's latency/error metrics, then its deployment history, then the error rate of the specific new version, then search logs for that version.

2. Database Saturation

Recognize it by: latency increase with normal CPU, no recent deployment, and a connection pool or query-latency metric near its limit.

Typical path: check the service's latency metric, confirm no recent deployment, then check the downstream database's connection-pool and query-latency metrics directly.

3. Memory Leak

Recognize it by: pod memory climbing continuously over hours (not a sudden spike), OOMKilled events, and rising restart counts.

Typical path: check Kubernetes pod status for restarts/OOMKilled events, then the memory metric trend over several hours.

4. Downstream Dependency

Recognize it by: the affected service's own CPU/memory is normal, but its latency tracks a dependency's latency.

Typical path: check the service's dependency graph, then check the latency metric of each dependency in turn.

5. Traffic Spike

Recognize it by: requests/sec, CPU, and latency all rising together, with no recent deployment.

Typical path: check requests/sec, confirm no recent deployment, then check whether autoscaling kept pace.

6. Diagnostic Tool Quick Reference

Use this table to decide which tool to call first based on the strongest available signal.

Tool	When to Use	Key Signal
get_metric	Any latency/error/resource question	current vs. baseline deviation

Tool	When to Use	Key Signal
get_deployment_history	Sudden change with unknown cause	a deploy shortly before the change
get_kubernetes_status	Suspected pod-level issue	restarts, OOMKilled, CrashLoopBackOff
get_service_dependencies	Service healthy but slow	map of upstream/downstream services
search_logs	Need root-cause detail after narrowing	error/timeout patterns in log text
get_recent_alerts	Establishing incident context/history	related alerts already fired